

RISC-V Microprocessor & DSP Element

Team Sachima

Ashwin K. Avula, Yucheng Chang, Garry Chen, Alvin Cheng

ECE 554: Microarchitecture Specification

University of Wisconsin-Madison

Table of Contents

Hardware Specification	3
Top-Level Interface	3
UART Bootloader Architecture	4
UART Protocol	5
RISC-V CPU Architecture	5
FETCH Stage (IF)	7
DECODE Stage (ID)	8
EXECUTE Stage (EX)	10
MEMORY Stage (MEM)	12
WRITE-BACK Stage (WB)	14
Data Hazard Management: Forwarding Unit	14
Control Hazard Management: Hazard Unit	15
DSP Accelerator Architecture	15
Basic Workflow	15
Image Buffer	16
Direct Memory Access Convention	17
Processing Element	19
Video Driver Architecture	21
Software Specification	22

Hardware Specification

Top-Level Interface

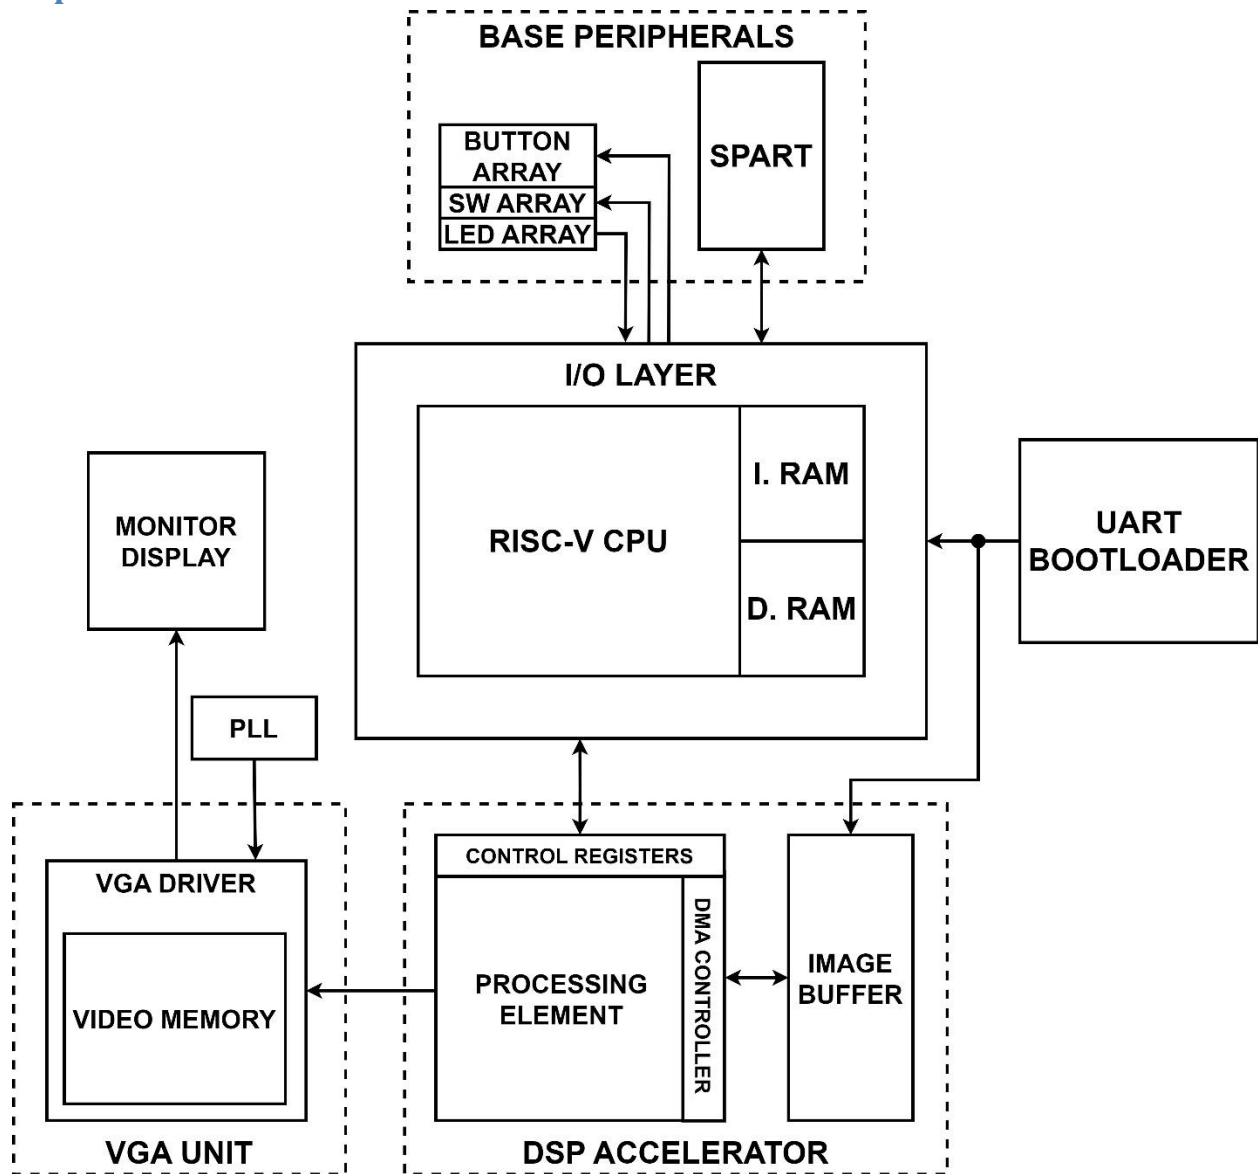


Figure 1: Top-Level Device Components

Depicted in *Figure 1*, the proposed device is controlled by a HOST RISC-V CPU core. Using an I/O Layer, the HOST can communicate with an array of peripherals via memory-mapped I/O. For example, a Special Purpose Asynchronous Receiver Transmitter (SPART) peripheral is used as a UART interface primarily for debugging purposes. The main DSP Unit employs a coprocessor architecture for accelerating various image processing operations. The DSP Unit interfaces with its own dedicated Image Buffer, which stores the 256x256 RGB images that the unit operates on. The unit utilizes Direct Memory Access (DMA) structures to maintain simultaneous read and write streaming of pixel data, in parallel with the image processing performed in the Processing Element of the DSP Unit. A dedicated video memory is employed to store the immediate images, which are output to a monitor display via a VGA

Driver. Finally, a UART Bootloading system is utilized for efficient reprogramming of the CPU Instruction and Data Memories, as well as the DSP's Image Buffer. The following sections will highlight interfaces between these primary components and the inner microarchitecture of each component. It is important to note that all modules are currently expected to map to the Altera SoC FPGA, and following discussions on memory modules detail the size constraints in order to meet this full utilization of on-chip RAM.

UART Bootloader Architecture

Depicted in *Figure 2*, a Bootloader system is utilized for the initialization of on-chip SRAM to avoid frequent and redundant compilation and flashing onto the Altera FPGA. Using a HOST PC, application code, image data, or raw data values can be converted to Binary and sent via UART over a COM Port in real-time. The main UART_boot module implements a general-purpose UART RX and, based on the protocol discussed in the following subsection, will store the incoming data into either the Instruction RAM, Data RAM, or the Image Buffer. During this bootloading process, the UART_boot acts as the master for the CPU and keeps the CPU in reset until bootloading is complete. Thus, execution will resume from PC Address 0x0000 when complete.

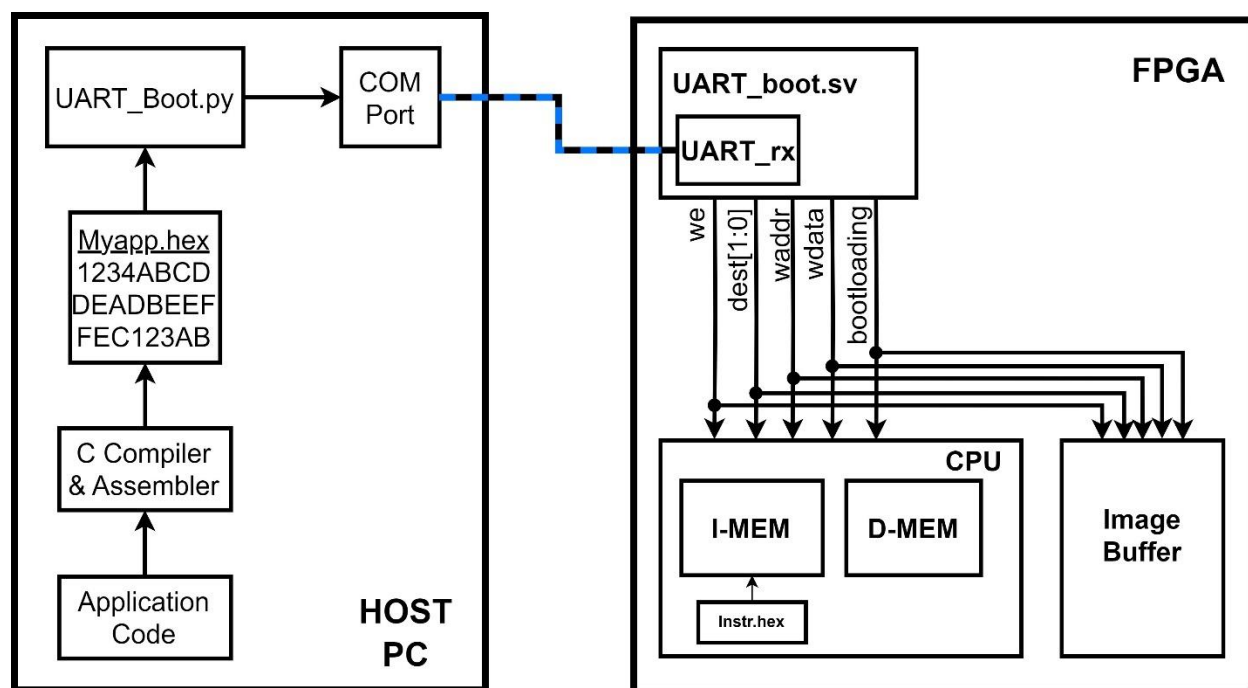


Figure 2: Top-Level Bootloader Architecture

Signal	Port Direction	Description
wdata[7:0]	Output	Serial Data Byte Output
we	Output	Active High Write Enable
done	Output	Transaction Complete Flag, Active High
Status[1:0]	Output	Status of the state machine: 0=IDLE, 1=CTRL, 2=ADDR, 3=DATA

bootloading	Output	Flag signifying Bootloading in process. Used to pull down Active Low CPU Reset
dst[1:0]	Output	Bootloaded Data Destination: 00=IM, 01=DM, 10=IB
RX	Input	Serial Bit Received from HOST PC

Table 1: UART Bootloader Interface

UART Protocol

The Bootloader module operates on a standard UART 8N1 frame with a baud rate of 115200 bps. The protocol requires the HOST CPU to initiate the transaction with a control byte, with the 2 least-significant bits indicating which memory is being written to. The next two bytes indicate the address of the memory to be written to. Finally, the following four bytes indicate the actual desired data to be written (assuming RISC-V 32b data width). This process repeats until all data is written to the desired memory, and the HOST CPU releases the UART Line high, completing the Bootloading process.



Figure 2: UART Bootloader Protocol

RISC-V CPU Architecture

For the main CPU core, the RISC-V Base 32I ISA is adopted. This implementation will include a pipelined architecture, a Harvard-style memory system, a 32 x 32-bit Register File, byte-addressable data memory, memory-mapped I/O, and compatibility with the open-source RISC-V toolchain. The following subsections detail each pipeline stage, memory constraints, and optimizations to minimize stalling and flushing occurrences.

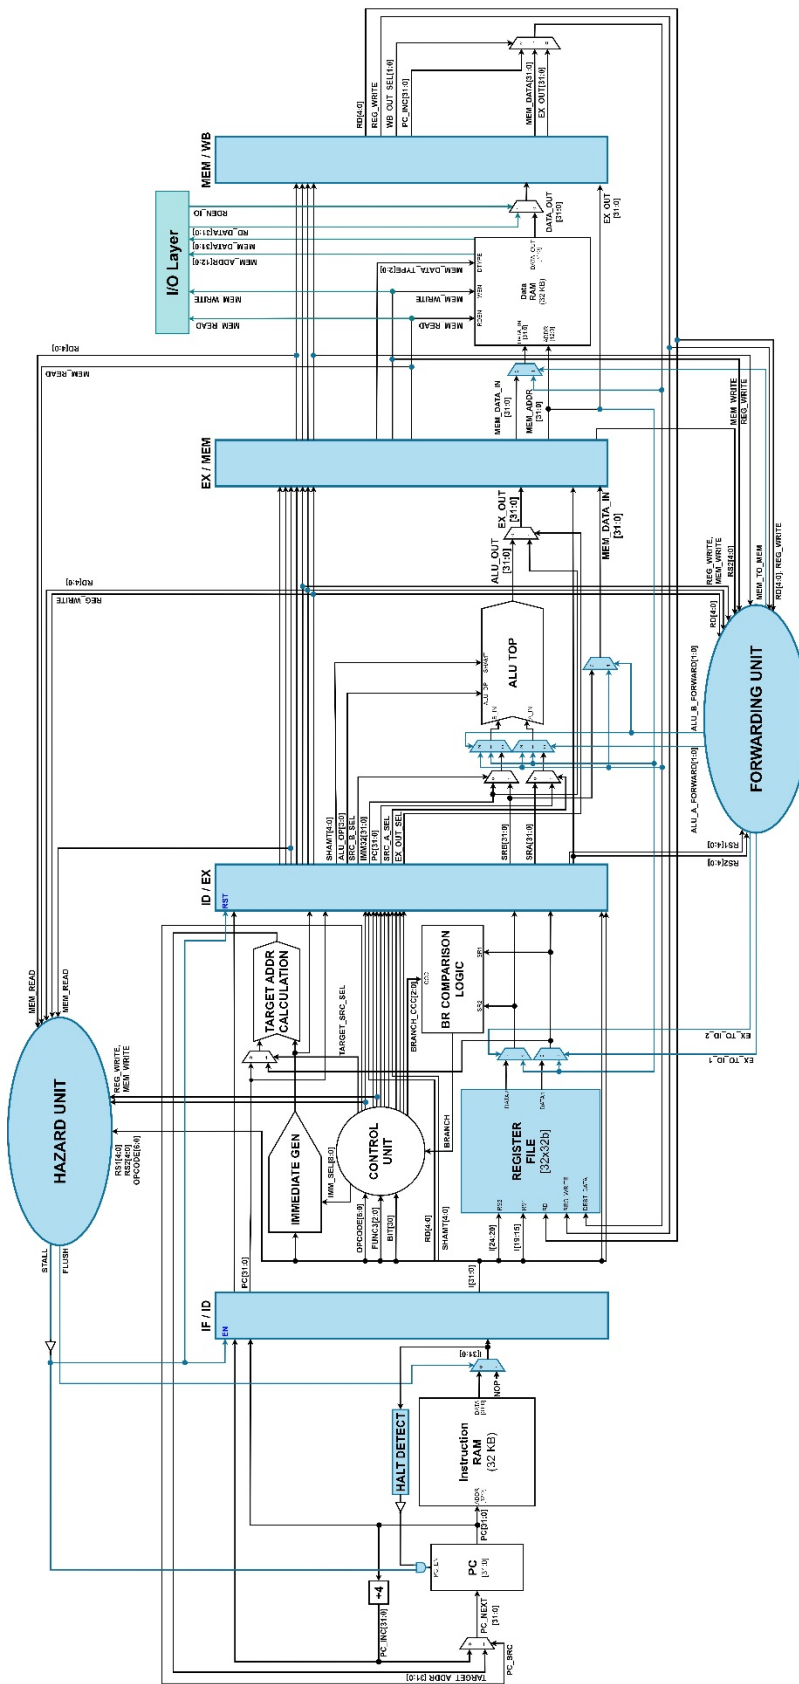


Figure 3: RISC-V Top-Level Pipeline Architecture

FETCH Stage (IF)

The Fetch pipeline stage consists of the PC register, which maintains the currently fetched instruction from Memory. In the event that the currently fetched instruction is an ECALL operation, the HALT DETECT Unit will signal a HLT Signal, which will disable the PC. Thus, the CPU will be stalled in this HALT state forever. The PC source mux selects between the incremented PC and the computed TARGET_ADDR in the event of control flow changes. The Fetch stage includes a 32KB Instruction Memory, which is a single-port single-cycle RAM. This memory was constructed to be a single-bank word-accessible memory, as instructions are only accessed in word granularity. Under the constraints of the FPGA on-chip RAM, the CPU was given only 64KB of memory to be used as instruction and data memories. Thus, the Instruction occupies 32 KB of this memory.

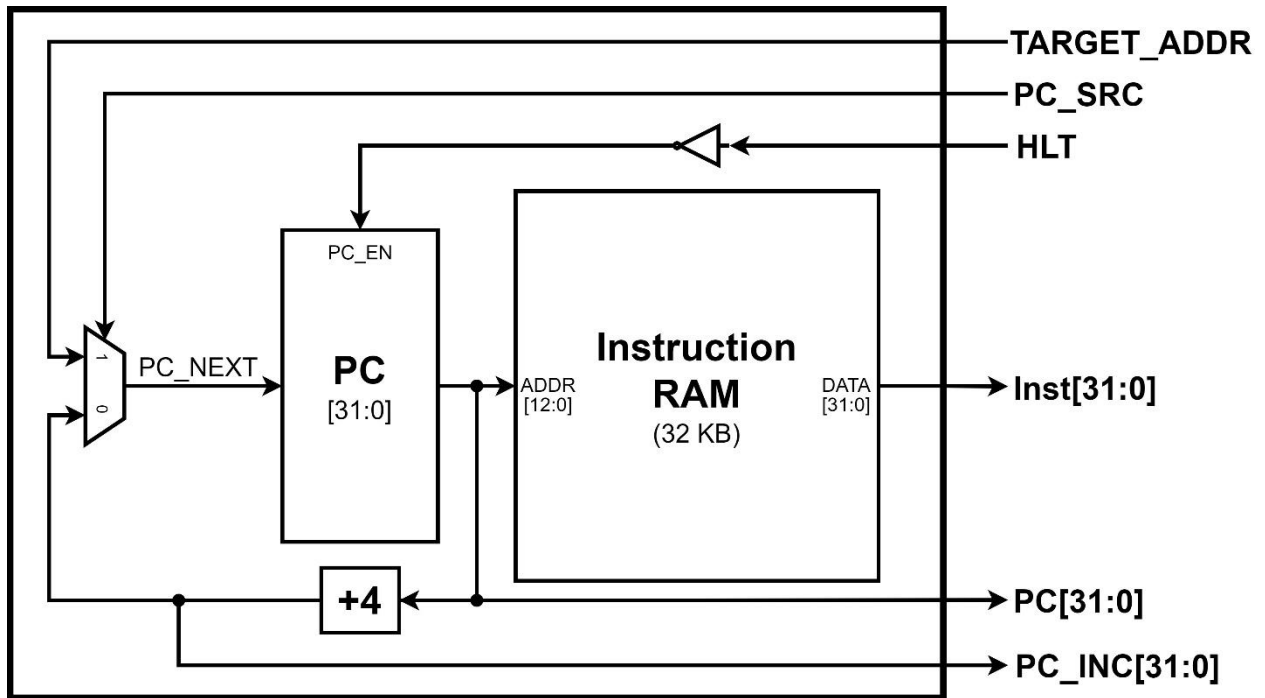


Figure 4: Top-Level Fetch Stage

Signal	Port Direction	Description	Origin Stage
TARGET_ADDR[31:0]	Input	32-Bit Address for Jump & Branch Cases	DECODE
PC_SRC	Input	PC_NEXT Mux Select	DECODE
HLT	Input	HALT Instruction Flag to disable PC	DECODE
PC[31:0]	Output	Current Program Counter	FETCH
PC_INC[31:0]	Output	Current Incremented Program Counter	FETCH
INST[31:0]	Output	32-Bit Instruction Fetched from Memory	FETCH

Table 2: FETCH Top Interface

DECODE Stage (ID)

The Decode pipeline stage handles the bulk of control generation, operand accessing, and branch/jump resolution. In this stage, the Control Unit generates all control signals using only the Instruction 7-bit Opcode, 3-bit Function field, BIT30, as well as a BRANCH flag from the BR Comparison Unit. With the generated control signals, the proper 32-Bit immediate can be generated, the correct branch/jump target address can be calculated, and the correct register locations can be accessed. All control signals not consumed in the current DECODE stage are passed to downstream stages and will propagate via pipeline registers.

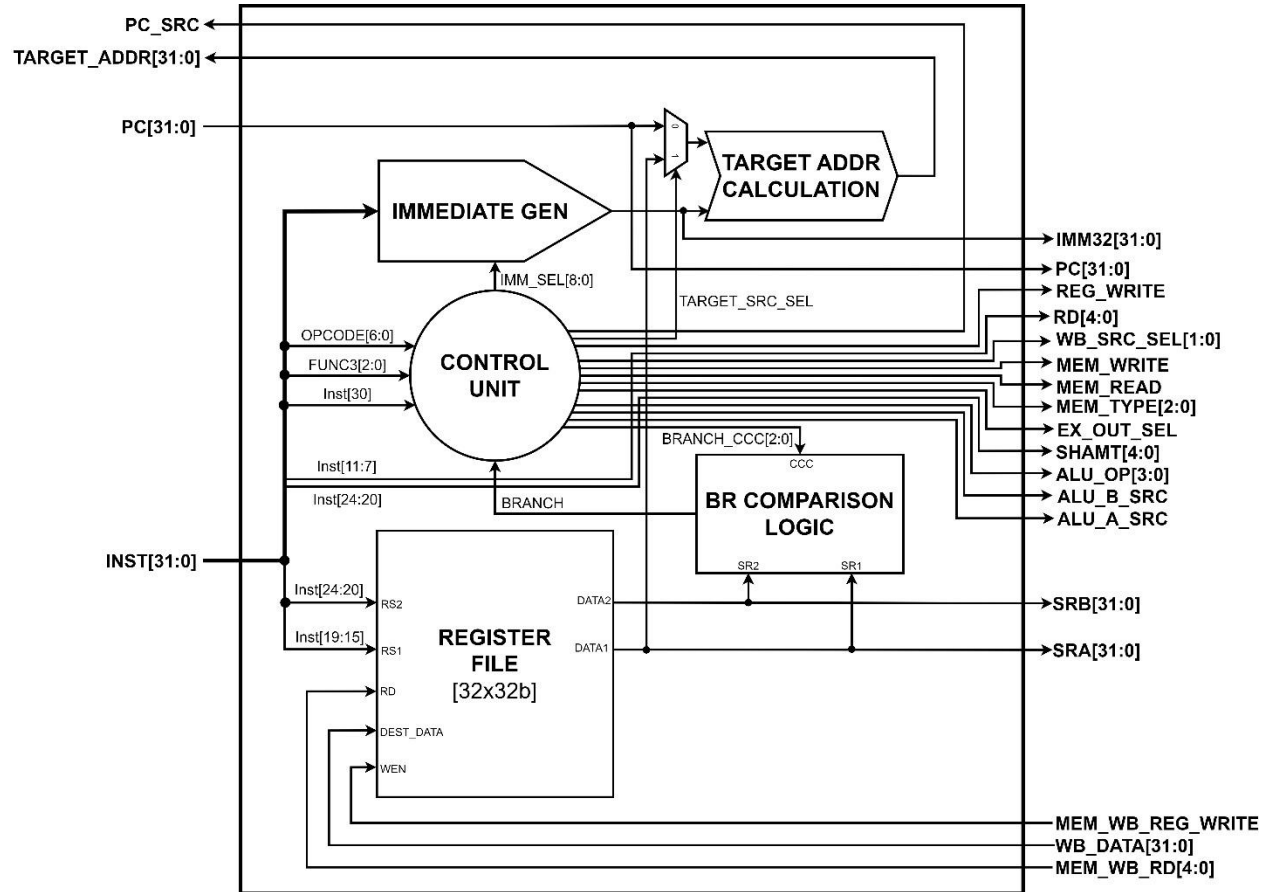


Figure 5: Top-Level Decode Stage

Signal	Port Direction	Description	Origin Stage
TARGET_ADDR[31:0]	Output	32-Bit Address for Jump & Branch Cases	DECODE
PC_SRC	Output	PC NEXT Mux Select	DECODE
HLT	Output	HALT Instruction Flag to disable PC	DECODE
PC[31:0]	Input, Output	Current Program Counter	FETCH
INST[31:0]	Output	32-Bit Instruction Fetched from Memory	FETCH
IMM32[31:0]	Output	32-Bit Immediate	DECODE

REG_WRITE	Output	Register File Write Enable	DECODE
RD[4:0]	Output	Register Final Destination Address	DECODE
WB_SRC_SEL[1:0]	Output	Write Back Stage Mux Select	DECODE
MEM_WRITE	Output	Data Memory Write Enable	DECODE
MEM_READ	Output	Data Memory Read Enable	DECODE
MEM_TYPE[2:0]	Output	Data Memory Access Type	DECODE
EX_OUT_SEL	Output	Execute Stage Output Mux Select	DECODE
SHAMT[4:0]	Output	ALU Shift Amount	DECODE
ALU_OP[3:0]	Output	ALU Control Function	DECODE
ALU_A_SRC	Output	ALU Operand A Mux Select	DECODE
ALU_B_SRC	Output	ALU Operand B Mux Select	DECODE
SRA[31:0]	Output	Register File Read Port 1 Output	DECODE
SRB[31:0]	Output	Register File Read Port 2 Output	DECODE
MEM_WB_REG_WRITE	Input	Pipelined Register File Write Enable	WRITE BACK
WB_DATA[31:0]	Input	Write Back Data Output, Register File Destination Data	WRITE BACK
MEM_WB_RD[4:0]	Input	Register File Destination Address	WRITE BACK

Table 3: DECODE Top Interface

Immediate Generation Unit




Figure 6: Immediate Generation Microarchitecture

Above, Figure 6 details the microarchitecture for the generation of Immediate Values. This unit follows the RISC-V Base 32I Immediate format and utilizes the ISA Instruction field sharing optimization. This minimizes the amount of selection logic, muxes, and wiring needed to generate the various types of immediates.

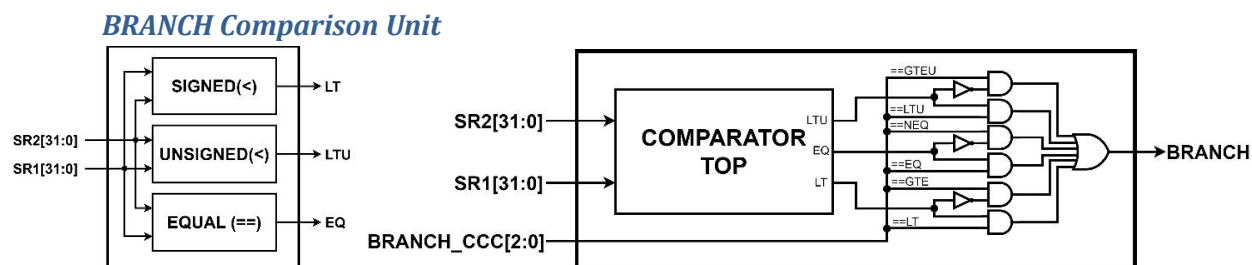


Figure 7: (A) General Purpose Comparator. (B) Branch Comparison Logic

Above, Figures 7A & B detail the microarchitecture for a general-purpose comparator that generates equality and inequality flags. This comparator is then implemented in the Branch Boolean Unit, which compares the Register File Read 1 & Read 2 Data outputs and, based on the instruction condition codes, determines if a branch should be taken. The comparator was designed to be general-purpose for implementation in the ALU as well.

EXECUTE Stage (EX)

The Execute pipeline stage consists of the main ALU, which handles the bulk of the computation for instructions in the RISC-V Base 32I ISA. The ALU selects between the Register File Data outputs, PC, and 32-Bit Immediate using its SRC muxes, and additional operation control generated in the Control Unit is used to select the instruction executed within the ALU. Finally, the ALU result is muxed with the 32-Bit Immediate in the event of an LUI operation, allowing the immediate to bypass the ALU and minimizing the control within the ALU.

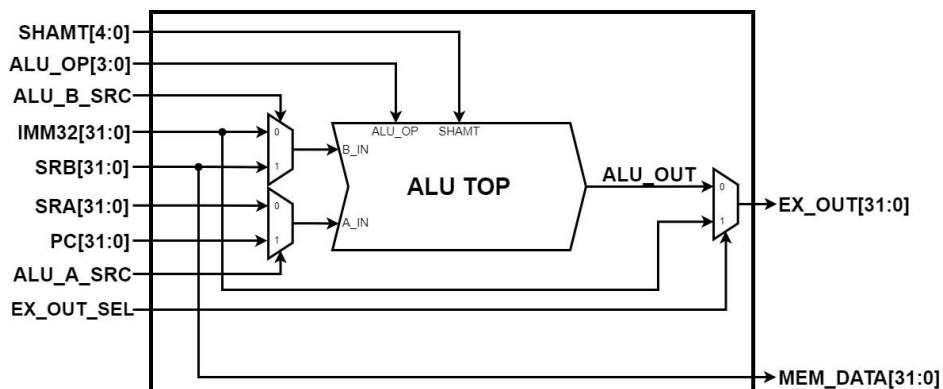


Figure 8: Top-Level EXECUTE Stage

Signal	Port Direction	Description	Origin Stage
IMM32 [31:0]	Input	32-Bit Immediate for I-type Instructions and address computation	DECODE

SRA[31:0]	Input	Register Data 1 Output for all ALU Instructions BUT AUIPC	DECODE
SRB[31:0]	Input	Register Data 2 Output for R-type Instructions	DECODE
PC[31:0]	Input	Program Counter for AUIPC	FETCH
SHAMT[4:0]	Input	Shift Amount for ALU Shift Instructions	DECODE
ALU_OP[3:0]	Input	ALU Control Vector	DECODE
ALU_A_SRC	Input	Mux Select for ALU A Operand	DECODE
ALU_B_SRC	Input	Mux Select for ALU B Operand	DECODE
EX_OUT_SEL	Input	Mux Select for Execute Output	DECODE
EXECUTE_OUT[31:0]	Output	Execute Output Result	EXECUTE

Table 4: EXECUTE Top Interface

Arithmetic-Logic Unit

The following *Figure x* details the microarchitecture of the primary ALU. For the main ADDER, a 32-bit Carry Lookahead structure is implemented for ADD, SUB, AUIPC, and LOAD/STORE instructions, as well as their immediate variants. A 3-function Barrel Shifter is implemented with a maximum shift of 31 bits for SLL, SRL, and SRA instructions, as well as their immediate variants. A simple 3-function Logical Unit is implemented for XOR, AND, and OR instructions, as well as their immediate variants. Finally, the Comparator detailed in the DECODE Section is reused in the ALU for SLT, SLTU instructions, as well as their immediate variants. All intermediate outputs are muxed to the ALU_OUT based on the ALU_OP.




Figure 9: BASE 32I ALU

MEMORY Stage (MEM)

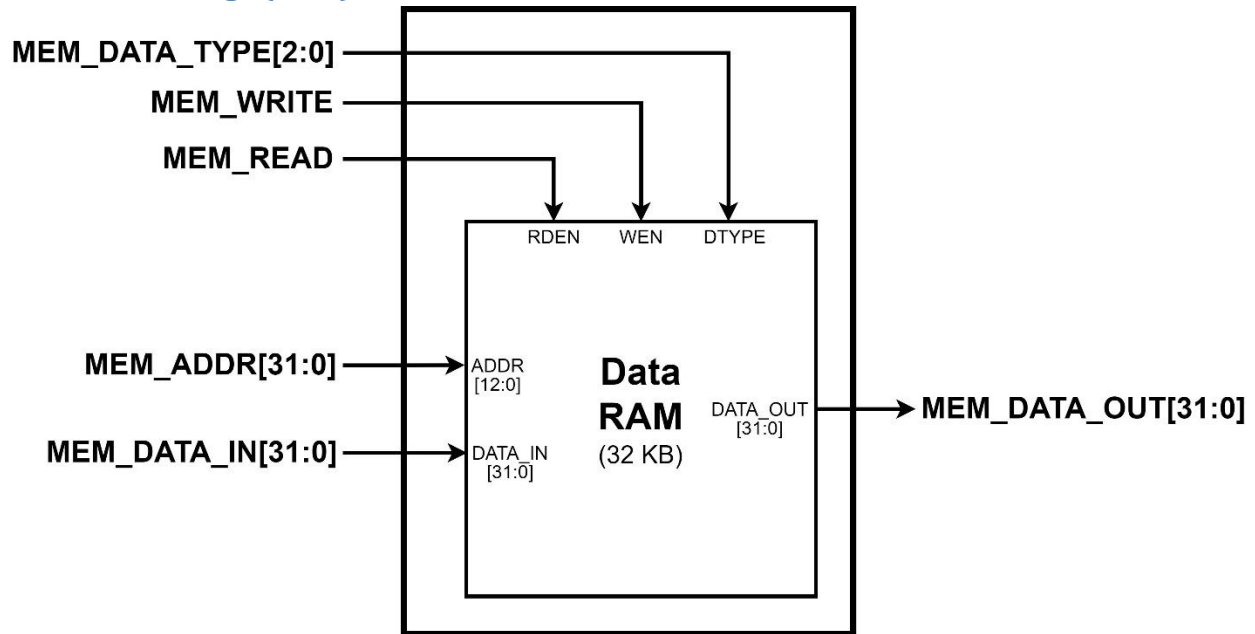


Figure 10: Top-Level EXECUTE Stage

Signal	Port Direction	Description	Origin Stage
MEM_DATA_TYPE[2:0]	Input	Data Memory Access Type	DECODE
MEM_WRITE	Input	Data Memory Write Enable	DECODE
MEM_READ	Input	Data Memory Read Enable	DECODE
MEM_ADDR[31:0]	Input	Byte-Address of Data Accessed in Current Cycle	EXECUTE
MEM_DATA_IN[31:0]	Input	Data to be Written on a Store Instruction	EXECUTE
MEM_DATA_OUT[31:0]	Output	Data Read from Memory to be Written to Register File on a Load Instruction	Memory

Table 5: MEMORY Top Interface

The Memory pipeline stage serves as the primary access point to Data Memory. This stage includes a 32KB Data Memory, which is a single-port single-cycle RAM. The memory was designed to be a four-bank byte-accessible memory, as data can be accessed in byte granularity (Store Byte, Load Byte, etc.). Under the constraints of the FPGA on-chip RAM, the CPU was allocated only 64KB of memory to be used as instruction and data memories. Thus, the Data occupies 32 KB of this memory.

In addition to accessing Data Memory, the Memory Stage allows for access to Memory-Mapped I/O. Note that Memory-Mapped access is omitted from Figure 10; however, the microarchitecture can be observed in *Figure 3*. Memory-Mapped Registers in the I/O layer can be found in *Table 6* below.

Memory Mapped Address	Function	CPU I/O Direction
0xC000	LED Array [9:0]	Output
0xC001	SW Array [9:0]	Input
0xC002	Button Array [3:0]	Input
0xC004	SPART Transmit Buffer [7:0]	Input / Output
0xC005	SPART Status Register [7:0]	Input
0xC006	SPART Division Buffer Low Byte [7:0]	Output
0xC007	SPART Division Buffer Low Byte [7:0]	Output
0xC00C	Coprocessor Status Register [7:0]	Input
0xC00D	Coprocessor Control Register [7:0]	Output

Table 6: Memory-Mapped I/O Region

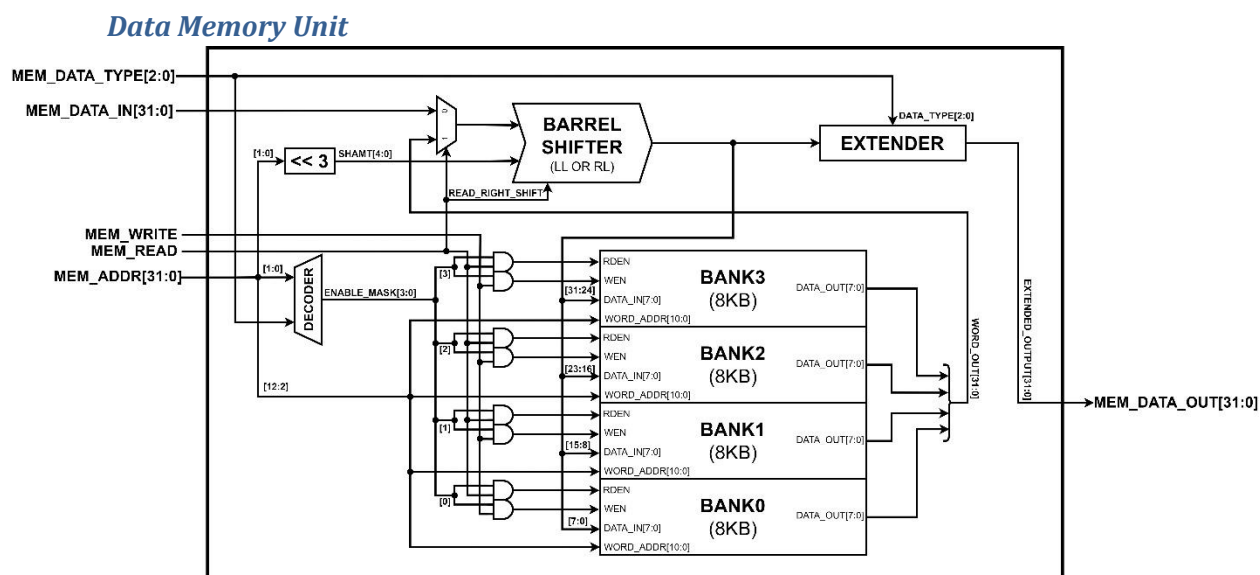


Figure 11: Data Memory Byte Banking Scheme

To achieve byte-addressability, the 32-bit Data Memory is banked into byte-regions. As detailed in the RISC-V Base 32I ISA, memory is assumed to be aligned with respect to the memory access data type. Furthermore, data input and output always begin at the least-significant byte. Thus, to input and output data from the various banks, a Right-or-Left Barrel shifter is implemented. By utilizing the least two significant bits of the memory address (Byte-Address), a bit address of the desired byte is obtained by left-shifting the byte address by 3 bits. This bit address is then used as a shift amount to right-shift Read data down to the least significant bit or to left-shift Write data up to the desired Byte Bank. Additionally, the output data is passed to an Extender unit, which manages zero and sign extension of non-word reads based on the memory data type. For example, given a Load-Byte instruction with an address having the least significant two bits of 10, a bit address of 16 is derived (reading from Byte Bank

2). Thus, the Data memory output is right-shifted so that the output of Byte Bank 2 now resides in the least-significant bit position.

WRITE-BACK Stage (WB)

The Write-Back pipeline stage simply selects the appropriate data value to be written into the Register File. This selection is made based on the WB_SRC_SEL control signal defined in the Control Unit. Additionally, Register File address and write enable signals are propagated to the Write-Back stage. This behavior is depicted in *Figure 12* below, and the interface is detailed in *Table 7*.




Figure 12: RISC-V WRITE-BACK Stage Top

Signal	Port Direction	Description	Origin Stage
WB_SRC_SEL[1:0]	Input	Mux Select control determining value passed to Reg File Dest Data	DECODE
PC_INC[31:0]	Input	Incremented PC propagated for JALx Instructions.	FETCH
MEM_DATA_OUT[31:0]	Input	Output of Data Memory for LOAD Instructions	MEMORY
EXECUTE_OUT[31:0]	Input	Output for all other instructions writing to the Reg File	EXECUTE
WB_DATA[31:0]	Output	Selected value to be written in the Reg File	WRITE-BACK
MEM_WB_RD[4:0]	Pass Through	Destination Address for value written to Reg File	DECODE
MEM_WB_REG_WRITE	Pass Through	Register File Write Enable	DECODE

Table 7: RISC-V WRITE-BACK Stage Interface

Data Hazard Management: Forwarding Unit

To handle Data Hazards introduced by the 5-stage pipeline, various forwarding methods are implemented to bypass pipeline registers and minimize the number of stall cycles required to retrieve the proper data for an instruction. This implementation utilizes: Register File Bypassing (Mem-to-

Decode) Forwarding, Ex-to-Decode Forwarding for Branch and Jump operands, Ex-to-Ex Forwarding and Mem-to-Ex Forwarding for ALU Operands, and finally Mem-to-Mem Forwarding for Store Instructions. The Forwarding Unit is an observing entity not bound to a single pipeline stage but instead spies on the control signals and addressing flowing through the pipeline to determine if forwarding is appropriate.

Control Hazard Management: Hazard Unit

Similar to the Forwarding Unit, the Hazard Unit is an observing entity for the management of Control Hazards introduced by the 5-stage pipeline. The Hazard Unit primarily introduces pipeline bubbles into the IF/ID registers in the event of a Jump or Branch taken, such that the currently fetched instruction is invalid and must be flushed. In the event that stalling is necessary, such as in load-to-use stall, data-to-branch, or data-to-jumpR cases, the FETCH & DECODE stages are frozen by disabling the IF/ID registers and resetting the ID/EX registers. This reset of the ID/EX registers introduces a pipeline bubble into the EXECUTE stage as instructions continue downstream to resolve the data hazard. More

DSP Accelerator Architecture

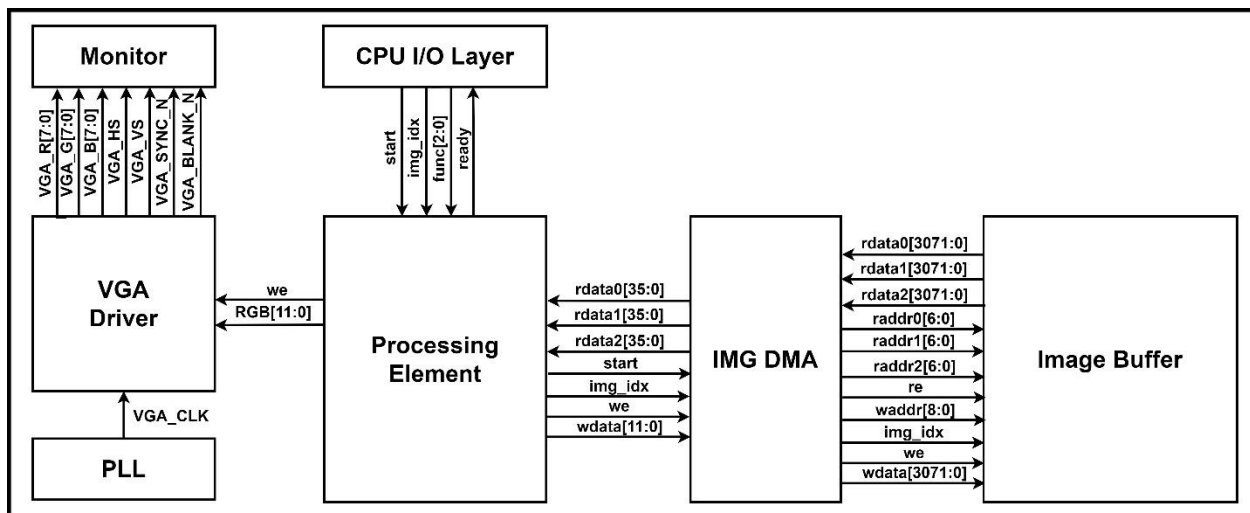


Figure 13: Top-Level Accelerator Architecture

The proposed device implements a coprocessor in the form of a DSP Accelerator to optimize various image pre-processing and processing functions. The DSP Accelerator consists of an Image Buffer for storage of an “Original” 256x256 input image and a “Most Recent” 256x256 output image. These two images are streamed to a Processing Element via the IMG DMA. The Processing Element is capable of executing various 3x3 Convolution filters on the streaming data as well as various pre-processing functions. Additionally, a VGA driver is implemented to output the processed image data to a Monitor Display. All data is secured to the coprocessor such that the HOST CPU can only initiate operations via memory-mapped I/O. The following sections detail the microarchitecture of these coprocessor components and the dataflow within the system.

Basic Workflow

On the very first startup, both the image buffer and the video memory will be empty. Each can store two images. The first image in both will contain the original image, while the second image will be the

most recently processed one. The UART bootloader will then load the original image into the image buffer, and the CPU will send instructions to the coprocessor to perform a NOP. This operation will compute various heuristic metadata of the image and then store the original image to the video memory for VGA to display. After this setup, the coprocessor will raise the ready signal when the operation is done, and all pixels streamed from the Image Buffer will be written into the video memory for display to the Monitor. Then, the CPU can continue to send functions and control bits to the coprocessor, and these operations will all be performed on the most recently processed image (functions can stack up).

Image Buffer

The image buffer stores two images: the original and the latest processed version. Each image has dimensions of 256x256 pixels, and every pixel encompasses 3 RGB channels, with each channel encoded with 4 bits. Furthermore, the images are segmented into three banks to support 3x3 kernel convolution operations efficiently. When reading or writing, the operations involve entire rows of 256 pixels (3072 bits) for each bank.




Figure 14: Image Buffer Banking Methodology

Signal	Port Direction	Description
img_idx	Input	Specifies the index of the image (2 images each for video memory and image buffer) for storing into different memory locations
raddr0 [6:0]	Input	Read address for bank0
raddr1 [6:0]	Input	Read address for bank1
raddr2 [6:0]	Input	Read address for bank2
re	Input	Read enable for all the banks in image buffer
we	Input	Write enable for the bank specified by waddr in image buffer
waddr [9:0]	Input	Bit 9 specifies the image index, bits 7-8 specifies the bank address, bits 0-6 specifies the bank address
wdata [3071:0]	Input	Data to be written to a row. 256 pixels * 4 pixels/channel * 3 channels
rdata0 [3071:0]	Output	Data to be read from bank0
rdata1 [3071:0]	Output	Data to be read from bank1
rdata2 [3071:0]	Output	Data to be read from bank2

Table 8: Coprocessor Image Buffer Interface

Image Buffer Memory

Memory	Description
Original Image [256x256x12]	The original image is loaded through the bootloader and consists of 256x256 pixels. This image is divided into three banks, with each bank containing 86 rows and columns of 256 pixels, amounting to 3072 bits per row. Notably, the final row in both bank 1 and bank 2 remains unoccupied, leaving it empty.
Recent Image [256x256x12]	The recent image, which is the most recent output from the processing unit, is also stored in the same format as the original image. However, the last rows in bank 1 and bank 2 of this recent image are repurposed as temporary storage areas. These rows are utilized to hold processed pixel rows.

Table 9: Image Buffer Size Specification

Direct Memory Access Convention

When the start signal is activated, the Direct Memory Access (DMA) retrieves data from the image buffer. It extracts three rows, each containing 256 pixels or 3072 bits, from the image buffer and outputs a 3x3 pixel window designated for convolution. Additionally, the DMA receives a processed pixel (12 bits) from the co-processor, stores it in the shifting register, and writes it back to the image buffer once the register is filled.

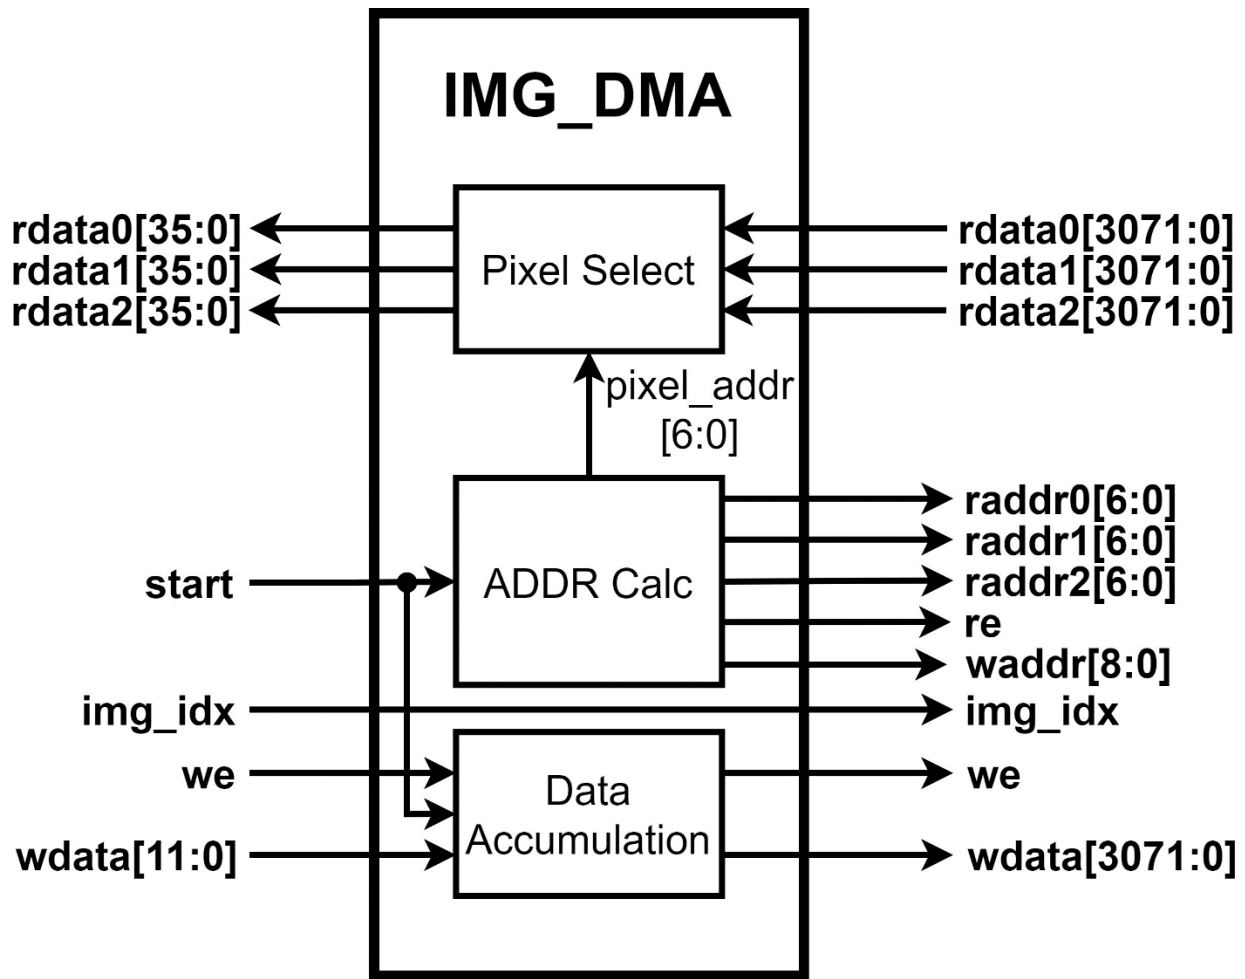


Figure 15: IMG DMA Top-Level

Signal	Port Direction	Description
rdata0_in [3071:0]	Input	Data read from bank0. Will be stored in registers
rdata1_in [3071:0]	Input	Data read from bank1. Will be stored in registers
rdata2_in [3071:0]	Input	Data read from bank2. Will be stored in registers
img_idx_in	Input	Specifies the index of the image (2 images each for video memory and image buffer) for storing into different memory locations
we_in	Input	Write enable input from coprocessor control unit
wdata_in [11:0]	Input	Processed pixel to be written. Will first be stored in a wide register to be written into image banks later
start	Input	Asserted for 1 clock to start memory access

rdata0_out [35:0]	Output	First row of data for a 3*3 convolution window, 3 pixels * 4 pixels/channel * 3 channels
rdata1_out [35:0]	Output	Second row of data for a 3*3 convolution window
rdata1_out [35:0]	Output	Third row of data for a 3*3 convolution window
img_idx_out	Output	Specifies the index of the image (2 images each for video memory and image buffer) for storing into different memory locations
raddr0 [6:0]	Output	Read address for bank0
raddr1 [6:0]	Output	Read address for bank1
raddr2 [6:0]	Output	Read address for bank2
re_out	Output	Read enable for all the banks in image buffer
we_out	Output	Write enable for the bank specified by waddr in image buffer
waddr [8:0]	Output	Bits 7-8 specifies the bank address, bits 0-6 specifies the bank address
wdata_out [3071:0]	Output	Data to write to a row. 256 pixels * 4 pixels/channel * 3 channels

Table 10: IMG DMA Interface

Direct Memory Access Register

Register	Description
Shifting Register [3071:0]	<p>The shifting registers served as the temporary storage of processed pixel from the processing unit. Each time the DMA received a processed pixel, it shifts the register right by 12 bit and store the processed pixel in the most significant 12 bits. When a row within the recent image is no longer required for current processing tasks, the Direct Memory Access (DMA) transfers the data stored in the shifting registers to that specific row.</p>

Table 11: IMG DMA Register Specification

Processing Element

The processing element acts as a dedicated image processor, executing tasks such as Gaussian Smoothing, sharpening, and edge detection, among others. Upon activation of the start signal, it streams a 3x3 pixel window from the DMA. If the grayscale bit is asserted by the control register set by the HOST CPU, the image undergoes conversion to grayscale before any processing. Following each processing step, the processor dispatches a single pixel (12 bits) to both the Direct Memory Access (DMA) for write-back to the Image Buffer and the Video Driver. Once the entire image has been processed, a ready signal is issued to indicate completion.

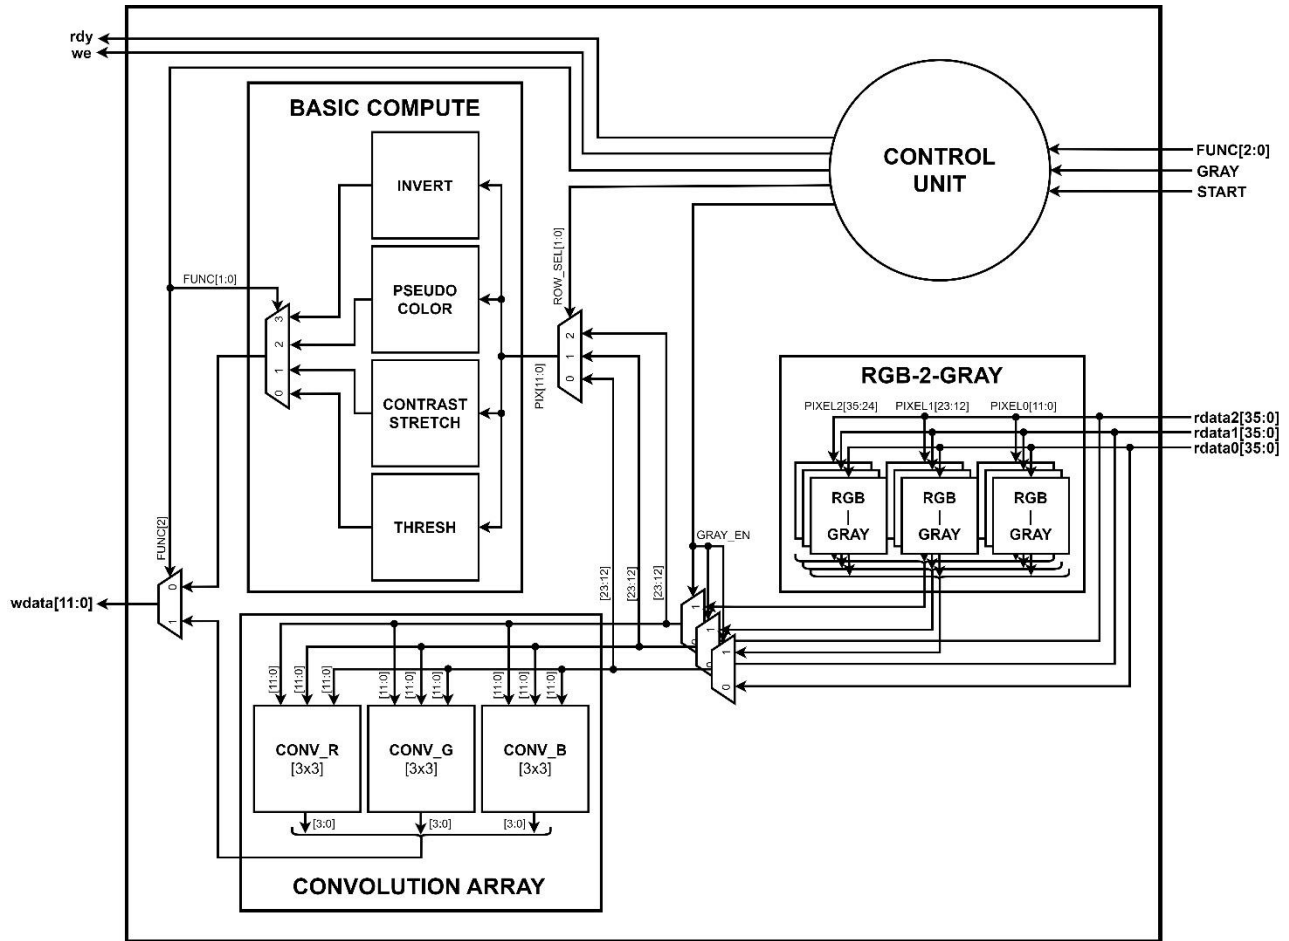


Figure 16: Processing Element Microarchitecture

Signal	Port Direction	Description
func [2:0]	Input	Indicates which image processing function will be performed. Function as opcodes in the coprocessor
grayscale	Input	Control bit to indicate if the image should be converted to grayscale before processing
img_idx	Input	Specifies the index of the image (2 images each for video memory and image buffer) for storing into different memory locations
rdata0 [35:0]	Input	Data read from bank 0 in image buffer
rdata1 [35:0]	Input	Data read from bank 1 in image buffer
rdata2 [35:0]	Input	Data read from bank 2 in image buffer
start	Input	Asserted for 1 clock to start coprocessor
RGB_out [11:0]	Output	Processed pixel output with four bits in each channel. If the image is grayscale, it will replicate 4 bits three times
we	Output	Write enable signal for both the image buffer and the video memory

ready	Output	Raise ready signal if the entire image is processed and written to video memory
-------	--------	---

Table 12: Processing Element Interface

Video Driver Architecture

The video driver architecture is designed to interface directly with the processing unit, from which it receives a 12-bit pixel after each processing operation. This pixel is first stored within the video memory. Following storage, the video driver expands the 12-bit pixel data to 24 bits. After the bit extension process, the 24-bit pixel is transmitted to the monitor for visualization. Finally, it is important to note that the VGA Monitor operates on a 25MHz clock, $\frac{1}{2}$ the rate of the device Refclk. Thus, a PLL is implemented for the 25MHz clock generation.

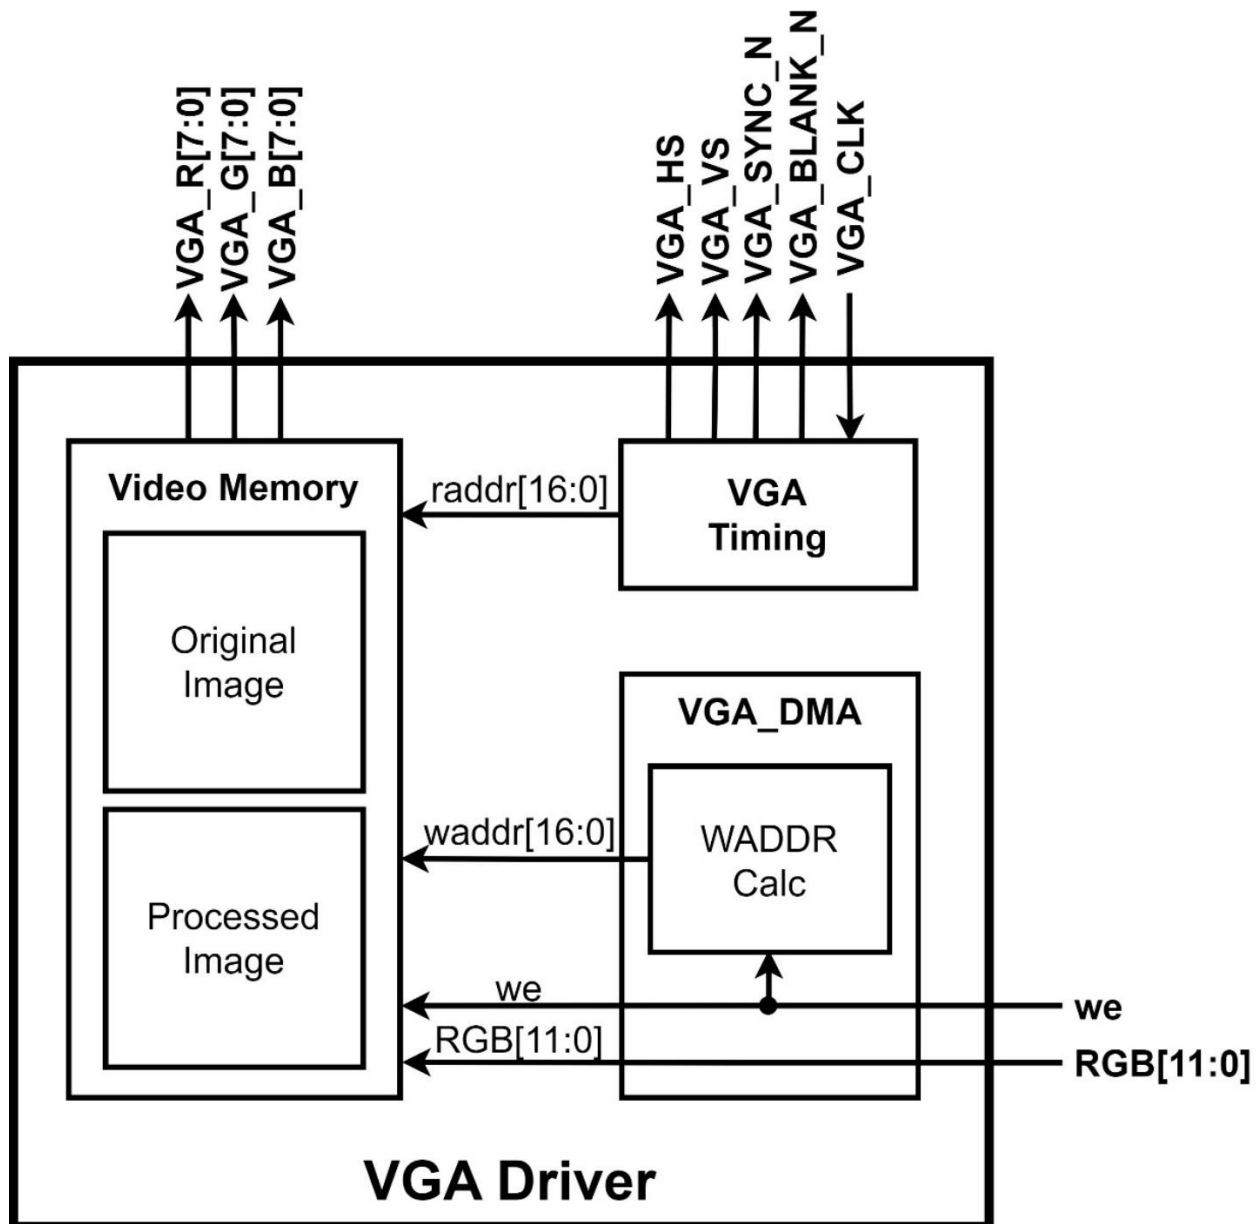


Figure 17: VGA Driver of Image Output to a Monitor Display

Signal	Port Direction	Description
we	Input	Write enable signal for video memory
RGB [11:0]	Input	Pixel data to be stored to video memory
VGA_HS	Output	Active low horizontal synch
VGA_VS	Output	Active low vertical synch
VGA_CLK	Output	25MHz VGA clock from PLL block
VGA_BLANK_N	Output	Assert (low) during non-active pixels
VGA_SYNC_N	Output	Tie it low
VGA_R [7:0]	Output	R channel, will be tied as {rdata[11:8],4'b0}
VGA_G [7:0]	Output	G channel, will be tied as {rdata[7:4],4'b0}
VGA_B [7:0]	Output	B channel, will be tied as {rdata[3:0],4'b0}

Table 13: VGA Driver Interface

Video Memory

Memory	Description
Images [640x256x12]	The original image is loaded to video memory at the start of application from the processing unit. Another image is the processed image that has been performed operations by the processing unit. Unlike the image buffer, the row of video memory stored the row of two different images since the images will be displayed side-by-side on the monitor.

Table 14: VGA Video Memory Specification

Software Specification

We will have a RISC-V processor and a customized coprocessor implemented on an FPGA to perform image processing tasks such as edge detection, highlighting, contrast adjustment, and recoloring. Our coprocessor is designed to execute image processing functions rapidly and efficiently, operating in parallel with the host CPU. The coprocessor supports fast and efficient parallel load and store functionalities from different memory banks and processes the image in filter batches. We will use switches on the FPGA board as inputs for selecting which filter to apply to the image. As small batches of the image are processed, the completed portions will be written to the video memory. Since the image data is read out of order in parallel, we need to reformat the processed data into the original image layout before writing to the video memory. Subsequently, the coprocessor will signal the CPU when all processing is complete, and the image is ready for the VGA to display. Both the original and processed images will be displayed side by side on the VGA monitor. The HOST CPU will primarily be utilized for polling peripherals and will behave as the control master for the coprocessor and the VGA driver.

We will compare the processing speed of the images (number of clock cycles used) between using the coprocessor and using the CPU solely. When using the CPU, the coprocessor will be detached, and the image data will be stored in the Data-Memory to suit the CPU.